

SemEval-2026-Task13-Subtask-A: Binary Machine-Generated Code Detection

Raviteja Gundu, Divya Navale, Divya Avuti

Abstract

The rapid progress of large language models (LLMs) has made it increasingly difficult to distinguish human-written code from machine-generated code. In this paper, we present our approach to SemEval-2026 Task 13 Subtask A ([mbzuai-nlp, 2025](#)), focusing on the binary classification of code authorship across seen and unseen programming languages. We evaluate a diverse set of architectures, including encoder-only (CodeBERT, GraphCodeBERT, ModernBERT), decoder-only (StarCoder), and unified transformers (UniXcoder), alongside feature-based baselines. Our experiments reveal that while large decoder-only models like StarCoder-3B achieve the highest overall performance on the original distribution, they are sensitive to training data volume. Conversely, encoder-based models benefit significantly from language-balanced training but struggle with cross-lingual generalization. Our findings suggest that while pretrained representations outperform classical features, achieving robustness under distribution shifts remains an open challenge.

1 Introduction

Recent advances in large language models (LLMs), such as GitHub Copilot and ChatGPT, have fundamentally transformed software development and computer science education. Trained on massive corpora of publicly available code, these models excel at tasks ranging from implementation to debugging, not only in widely used languages like Python and Java but also in less common ones. This ability to generalize across diverse programming languages and application contexts has lowered entry barriers for beginners and boosted productivity for experts, driving the rapid adoption of AI-based coding assistants in both academic and industrial settings.

However, the same capabilities that make LLMs powerful also introduce a critical challenge: dis-

tinguishing human-written code from machine-generated code has become increasingly difficult. As models improve at mimicking human coding styles across multiple programming languages and domains, the boundary between human and machine-generated code continues to blur. This ambiguity is not merely theoretical—it has tangible consequences for academic integrity, intellectual property protection, and software security.

This creates an urgent need for reliable detection mechanisms. While existing detection systems often achieve high accuracy on familiar programming languages and well-defined domains, such as algorithmic problem-solving, they frequently fail to generalize to unseen languages (e.g., JavaScript, Go, PHP). This lack of robustness under distribution shifts severely limits their real-world applicability.

In this work, we focus on the challenge of building a generalizable binary detector for machine-generated code. Specifically, we investigate the following research question: *How effectively can a detection model distinguish human-written from machine-generated code when evaluated on previously unseen programming languages?* Through systematic out-of-distribution evaluation, we quantify the generalization gap of existing approaches and explore strategies to improve detection robustness across linguistic variations.

This problem setting follows the definition of Subtask A of SemEval-2026 Task 13, *Detecting Machine-Generated Code* ([mbzuai-nlp, 2025](#)).

2 Related Work

The growing integration of large language models (LLMs) into software development has prompted significant research interest in distinguishing human-written from machine-generated code. Early efforts have primarily focused on binary classification of generated code in limited program-

ming languages and domains.

Nguyen et al. (Nguyen et al., 2024) introduced *GPTSniffer*, a CodeBERT-based binary classifier designed to detect ChatGPT-generated code in Python and Java using the CodeSearchNet dataset (Husain et al., 2019). Their approach achieved strong in-domain accuracy but struggled to generalize to unseen programming languages, highlighting the limitations of dataset diversity in code authorship detection. Similarly, (Idialu et al., 2024) proposed a stylometric-based detection framework that leveraged code-level writing style features to identify GPT-4-generated code at the class granularity. While effective for stylistic differentiation, such methods depend heavily on consistent structural and lexical features, reducing their robustness across varied code domains.

Orel et al. (Orel et al., 2024) introduced *Droid*, a comprehensive resource suite for machine-generated code detection covering multiple programming languages, generation models, and problem settings. Their work contributed two components relevant to this study: (1) a data filtering pipeline that removes low-quality, boilerplate, and near-duplicate samples to increase authorship signal strength, and (2) a CatBoost-based classification baseline using gradient-boosted decision trees over lexical and structural code features. These components demonstrated competitive performance under cross-model and cross-language evaluations, reinforcing the effectiveness of feature-driven detection methods in broader authorship attribution scenarios.

3 Methodology

We formulate SemEval-2026 Task 13 Subtask A as a binary classification problem. Given a code snippet x , the detector predicts a label $y \in \{0, 1\}$ where 0 denotes human-written code and 1 denotes machine-generated code. To study robustness under cross-language we evaluate transformer-based detectors spanning different architectural families, along with feature-driven classical baselines.

3.1 Transformer Families

To cover diverse inductive biases in code modeling, we evaluate three transformer families:

Encoder-only Transformers learn bidirectional contextual representations and are commonly used for classification via a pooled embedding. We fine-tune *CodeBERT*, *GraphCodeBERT*, and *Modern-*

BERT by adding a lightweight classification head on top of the pooled sequence representation. The resulting model is optimized end-to-end for binary prediction.

Decoder-only Transformers are autoregressive and are pretrained for next-token prediction, often capturing strong generative priors and long-range code patterns. We evaluate *StarCoder-3B* by adapting the decoder representations for sequence classification using a task-specific classification head.

Unified Transformers combine bidirectional encoding with autoregressive decoding capabilities, which can be beneficial for learning both global and local signals. We evaluate *UniXcoder* as a unified code model and fine-tune it for binary classification with a classification head over the sequence representation.

3.2 Parameter-Efficient Fine-Tuning

For large models, we employ parameter-efficient fine-tuning to reduce memory and compute overhead. In particular, we use *PEFT* with *LoRA* adapters for CodeBERT and StarCoder-3B, updating only a small set of trainable low-rank matrices while keeping the backbone largely frozen. This allows stable fine-tuning under constrained GPU budgets while preserving most pretrained knowledge.

3.3 Language Balancing Strategy

The official training split exhibits strong language skew toward Python, which may encourage detectors to overfit to Python-specific stylistic artifacts rather than general authorship cues. To mitigate this risk, we create a balanced training set by down-sampling Python to 50K total instances (stratified by label) while retaining all C++ and Java instances. A similar stratified downsampling strategy is applied to the validation split. Sampling is performed with a fixed random seed (random state = 42) to ensure reproducibility.

3.4 Feature-Based Baselines

To benchmark transformer detectors against classical machine-learning approaches, we implement two feature-driven baselines: *Random Forest* and *CatBoost*. Both models are trained on engineered lexical and structural features extracted from code snippets. These baselines provide a reference point for measuring the benefit of pretrained represen-

178 tations and help identify whether simple stylistic
179 features can generalize across languages.

180 4 Experimental Setup

181 4.1 Data Splits

182 We use the official SemEval-2026 Task 13 dataset
183 for Subtask A, restricted to human-written and
184 machine-generated samples. The original training
185 set contains 500K examples and the original validation
186 set contains 100K examples, with a strong
187 language skew toward Python. The test set is kept
188 unchanged (1K samples) to ensure comparability
189 across all experiments.

190 4.2 Training/Validation Distributions

191 The original training distribution is dominated by
192 Python (457,306 samples), whereas C++ and Java
193 together contribute fewer than 45K instances. To
194 reduce this imbalance, we construct a balanced
195 training split by downsampling Python to 50K total
196 while preserving all C++ and Java examples.
197 Similarly, we downsize validation from 100K to
198 10K using stratified sampling across language and
199 label.

Language	Label	Original	Balanced
C++	Human	11,147	11,147
	Machine	12,245	12,245
Java	Human	9,225	9,225
	Machine	10,077	10,077
Python	Human	218,103	23,760
	Machine	239,203	26,240
Total		500,000	92,694

200 Table 1: Training set distribution before and after
201 Python downsampling.

202 4.3 Tokenization and Input Formatting

203 All transformer models use their corresponding
204 AutoTokenizer. Inputs are padded/truncated to
205 model-dependent maximum sequence lengths. For
206 experiments where preprocessing is enabled (notably
207 for StarCoder variants), we additionally apply
208 lightweight normalization intended to stabilize
209 authorship signals (e.g., trimming overly long scaf-
210 folding and normalizing superficial literals where
applicable). These steps are treated as controlled
ablations rather than mandatory preprocessing.

211 4.4 Optimization and Training Protocol

212 All models are fine-tuned end-to-end (or via LoRA
213 adapters for PEFT runs) for a fixed duration of
214 3 epochs. We use the AdamW optimizer with a
215 learning rate of 2×10^{-5} and standard cross-entropy
216 loss for binary classification.

217 For experiments on the original dataset, we stan-
218 dardize the maximum sequence length to 512 to-
219 kens and use a global batch size of 256. For ex-
220 periments on the balanced dataset, we increase the
221 context window to capture longer-range depen-
222 dencies, utilizing maximum sequence lengths of 512
223 and 8192 depending on the model’s architecture.
224 Correspondingly, batch sizes for the balanced split
225 are adjusted per model to accommodate these larger
226 contexts within memory constraints. At inference,
227 predictions are obtained from the model logits us-
228 ing a fixed threshold of 0.5.

229 4.5 Evaluation Metrics

230 We evaluate using **accuracy**, and **macro F1**.
231 Macro-averaging ensures that the classes contribute
232 equally to the final score, which is important when
233 measuring generalization across languages.

234 5 Results

235 5.1 Overall Performance on the Original 236 Dataset

Model	Precision	Recall	Accuracy	Macro F1
StarCoder-3B	0.5720	0.6033	0.55	0.5230
UniXcoder	0.6058	0.5905	0.3860	0.3849
CodeBERT	0.5317	0.5208	0.31	0.304
ModernBERT	0.5965	0.5438	0.3010	0.2876
GraphCodeBERT	0.4453	0.4615	0.2750	0.2702
CatBoost	0.5696	0.5339	0.2980	0.2859
RandomForest	0.2440	0.9110	0.3500	0.3480

237 Table 2: Overall performance of models on the original
238 dataset.

239 As shown in Table 2, StarCoder-3B achieves the
240 strongest overall performance, with an accuracy of
241 0.55 and a macro F1 score of 0.5230. This demon-
242 strates the effectiveness of large decoder-only pre-
243 trained code models for machine-generated code
244 detection under the original, Python-dominated
245 training distribution.

246 Among encoder-only models, UniXcoder
247 achieves high macro precision (0.6058) and
248 recall (0.5905) but substantially lower accuracy
249 (0.3860), suggesting unstable decision boundaries
when generalizing across languages. CodeBERT,
ModernBERT, and GraphCodeBERT exhibit

250 weaker performance, with macro F1 scores below
251 0.31, highlighting the difficulty of general-purpose
252 encoders in this cross-domain setting.

253 Feature-based baselines show contrasting behavior.
254 CatBoost achieves moderate precision but low
255 overall accuracy, while Random Forest attains very
256 high recall (0.9110) at the expense of extremely
257 low precision (0.2440), indicating a strong bias to-
258 ward predicting machine-generated code. These
259 results confirm that classical models relying solely
260 on surface-level features struggle to balance predic-
261 tions under distribution shift.

262 5.2 Overall Performance on the Balanced 263 Dataset

Model	Accuracy	Macro F1
StarCoder-3B	0.2720	0.2699
UniXcoder	0.2820	0.2745
CodeBERT	0.2970	0.2914
ModernBERT	0.3310	0.3246
GraphCodeBERT	0.3920	0.3915
RandomForest	0.3500	0.3480

264 Table 3: Overall performance of models on the balanced
265 dataset.
266

267 Table 3 reports results for models trained on the
268 balanced dataset constructed as described in Sec-
269 tion 4.1. In this setting, GraphCodeBERT achieves
270 the best performance, reaching an accuracy of
271 0.3920 and a macro F1 score of 0.3915. This im-
272 provement suggests that structural representations
273 become more effective when dominant language
274 artifacts are reduced.

275 ModernBERT also benefits from balancing, im-
276 proving to 0.3310 accuracy and 0.3246 macro F1,
277 outperforming CodeBERT and UniXcoder. In con-
278 trast, StarCoder-3B experiences a notable drop in
279 performance when trained on the balanced dataset,
280 falling to 0.2720 accuracy and 0.2699 macro F1.
281 This indicates that aggressive downsampling may
282 remove useful training diversity for large decoder-
283 only models that benefit from larger-scale expo-
284 sure.

285 Overall, these results reveal a trade-off between
286 reducing language skew and preserving sufficient
287 training signal. Encoder-based and structure-aware
288 models benefit more consistently from balancing,
289 while decoder-only models achieve higher peak
290 performance on the original dataset.

291 5.3 Per-Language Performance: Seen vs. 292 Unseen Languages

293 Table 4 presents per-language macro F1 scores
294 for models trained on the original dataset. The
295 test languages are divided into seen (Java, C++,
296 Python) and unseen (C, C#, Go, JavaScript, PHP)
297 languages.

298 For seen languages, all models achieve their
299 highest and most stable performance, particularly
300 on Python. StarCoder-3B shows strong macro
301 F1 on Java and C++, while encoder-based mod-
302 els such as GraphCodeBERT and ModernBERT
303 remain comparatively weaker even on seen lan-
304 guages, indicating limited capacity to exploit train-
305 ing familiarity.

306 For unseen languages, performance varies sub-
307 stantially across models. StarCoder-3B demon-
308 strates strong generalization to Go (0.6738) and
309 competitive performance on C and C++, suggest-
310 ing that large decoder-only models capture trans-
311 ferable authorship cues beyond training languages.
312 Encoder-based models, particularly GraphCode-
313 BERT and ModernBERT, show significant degra-
314 dation on unseen languages such as PHP, where
315 macro F1 drops below 0.20.

316 Feature-based methods such as CatBoost per-
317 form poorly across most unseen languages, rein-
318 forcing the difficulty of relying on surface-level
319 features for cross-language authorship detection.
320 Notably, all models struggle on PHP, highlighting
321 it as a particularly challenging unseen language
322 under the current training regime.

323 5.4 Summary of Findings

324 Across all experiments, large pretrained code mod-
325 els consistently outperform classical baselines.
326 Decoder-only transformers achieve the highest
327 overall performance on the original dataset, while
328 encoder-based and structurally aware models show
329 improved robustness when training data is balanced.
330 The per-language analysis reveals substantial per-
331 formance gaps between seen and unseen languages,
332 underscoring the importance of multilingual pre-
333 training and diverse training distributions for robust
334 machine-generated code detection.

335 6 Additional Analysis

336 6.1 Exploratory Experiments with Larger 337 Models

338 We additionally explored a larger decoder-only
339 model, StarCoder-7B, to assess whether increased
340

Language	StarCoder-3B	UniXcoder	CatBoost	GraphCodeBERT	ModernBERT
C	0.4811	0.4491	0.1758	0.1911	0.2477
C#	0.4186	0.3306	0.2043	0.1586	0.1984
C++	0.4663	0.3599	0.2788	0.2532	0.2768
Go	0.6738	0.3667	0.2105	0.2734	0.2105
Java	0.4772	0.3434	0.1841	0.1957	0.1777
JavaScript	0.4118	0.4578	0.2944	0.1986	0.2735
PHP	0.4023	0.1652	0.1652	0.1037	0.0588
Python	0.4832	0.4429	0.4301	0.4309	0.4561

Table 4: Per-language Macro F1 scores for each model.

model capacity improves machine-generated code detection. While StarCoder-7B performs competitively, it does not consistently outperform the 3B variant, indicating diminishing returns under limited fine-tuning and compute budgets. Furthermore, StarCoder-7B exhibits more consistent behavior across unseen languages but does not achieve the same peak performance observed with the smaller model. These observations suggest that increased scale alone is insufficient to guarantee stronger generalization and that larger models may require more extensive optimization to fully realize their potential.

6.2 Preprocessing Strategies

We experimented with additional preprocessing strategies for StarCoder-3B, specifically applying: (1) insertion of language tags to preserve cross-lingual signals, (2) smart line-based trimming to remove boilerplate, and (3) normalization of literals to reduce superficial stylistic artifacts. Inputs were padded or truncated to a fixed sequence length. However, the preprocessed StarCoder-3B variant did not outperform the best baseline configuration.

Model	Precision	Recall	Accuracy	Macro F1
StarCoder-3B	0.5720	0.6033	0.55	0.5230
StarCoder-3B (preprocessed)	0.5411	0.5577	0.4940	0.4742
StarCoder-7B	0.5179	0.5258	0.5190	0.4784

Table 5: Exploratory analysis of model scale and preprocessing effects for StarCoder.

6.3 Limitations and Error Patterns

Across all architectures, PHP remains the most challenging unseen language, with consistently low macro F1 scores. This suggests that language-specific syntax, library usage, or domain differences introduce patterns not adequately captured by current detectors. Additionally, models frequently confuse human-written boilerplate with machine-generated code, indicating that stylistic similarity remains a key source of error.

7 Conclusion

In this work, we investigated the effectiveness of transformer-based models and classical baselines for detecting machine-generated code. Through systematic evaluation on SemEval-2026 Task 13 Subtask A, we demonstrated that architecture choice significantly influences robustness. Our results indicate that large decoder-only models, specifically StarCoder-3B, provide the strongest generalization capabilities on the original distribution, successfully transferring knowledge to some unseen languages like Go. However, we found that encoder-based models such as GraphCodeBERT and ModernBERT offer better stability when trained on balanced datasets, emphasizing the trade-off between model scale and data distribution strategies.

Despite these successes, a significant generalization gap remains. All evaluated models exhibited severe performance degradation on specific unseen languages, most notably PHP, suggesting that current detectors rely heavily on language-specific syntactic patterns rather than universal traits of machine-generation. Future work will focus on narrowing this gap by exploring cross-lingual contrastive alignment, incorporating more granular structural features, and extending detection capabilities to hybrid human-machine collaborative scenarios.

8 Github Link

The code and experimental results for this paper are available at: <https://github.com/ravitejagundu11/Polyglot-Code>.

References

- Hamel Husain, Ho-Hsiang Wu, Tiferet Gazit, Miltiadis Allamanis, and Marc Brockschmidt. 2019. Code-searchnet challenge: Evaluating the state of semantic code search. *arXiv preprint arXiv:1909.09436*.

408 Precious Idialu, Xiaolong Wang, and Md Tarek Rah-
409 man. 2024. Stylometric and structural analysis for
410 gpt-4 code authorship detection. In *Proceedings of*
411 *the 2024 AAAI Conference on Artificial Intelligence*.
412 AAAI Press.

413 mbzuai-nlp. 2025. Semeval-2026 task 13: Detecting
414 machine-generated code. <https://github.com/>
415 `mbzuai-nlp/SemEval-2026-Task13`. Accessed
416 October 2025.

417 Phuong T. Nguyen, Juri Di Rocco, Claudio Di Sipio,
418 Riccardo Rubei, Davide Di Ruscio, and Massimil-
419 iano Di Penta. 2024. Gptsniffer: A codebert-based
420 classifier to detect source code written by chatgpt.
421 *Journal of Systems and Software*, 214:112059.

422 Daniil Orel, Indraneil Paul, Iryna Gurevych, and Preslav
423 Nakov. 2024. Droid: A resource suite for AI-
424 generated code detection. In *Proceedings of the Con-*
425 *ference on Empirical Methods in Natural Language*
426 *Processing (EMNLP)*.